

If you know all these, then you know most things in GenAI

Prompt Engineering

Prompt Techniques

1. Chain of Thought (CoT)
2. Few-Shot Chain of Thought (Few-Shot-CoT)
3. ReAct (Reasoning and Acting)
4. Tree of Thoughts (ToT)
5. Self-Consistency
6. Hypothetical Document Embeddings (HyDE)
7. Least-to-Most Prompting
8. Graph Prompting
9. Recursive Prompting
10. Generated Knowledge
11. Automatic Reasoning and Tool-Use (ART)
12. Automatic Prompt Engineer (APE)

Transformer Architecture

1. Self-Attention Mechanism
2. Positional Encoding
3. Multi-Head Attention:
4. Encoder-Decoder Architecture
5. Layer Normalization and Feed-Forward Layers
6. Pre-training and Fine-tuning
7. Scalability with Parallelization
8. Applications Across Domains

Cloud Support

Azure – Azure OpenAI, Azure AI Studio, Azure Machine Learning

AWS – Amazon Bedrock, Amazon Sagemaker, Amazon Q

Google – Vertex AI

LLMOPs

1. Model Deployment and Scaling
2. Monitoring and Logging
3. Versioning and Model Lifecycle Management
4. Cost Optimization
5. Feedback Loops and CI
6. Compliance and Security
7. Latency and Throughput Optimization
8. A/B Testing

LLM Evaluation

1. BLEU (Bilingual Evaluation Understudy)
2. ROUGE (Recall-Oriented Understudy for Gisting Evaluation)
3. Perplexity
4. Exact Match (EM)
5. Human Evaluation Scores
6. Bias and Fairness Metrics
7. Toxicity Scores

LLM Frameworks (Across Layers) – LangChain/LlamaIndex

RAG

Vector DB - Pinecone, Weaviate, Qdrant, Chroma, Milvus, Vespa, LanceDB

Embedding Models – OpenAI's Embedding, Google's text-embedding-004, Huggingface Open-source Embedding Models

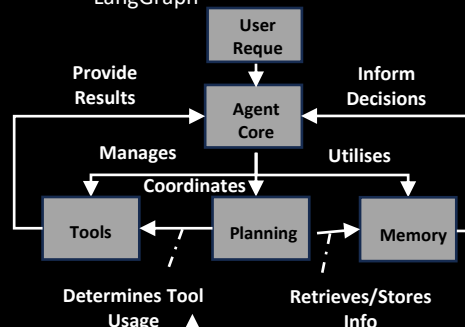
RAG Techniques

1. Basic RAG
2. Re-ranking RAG
3. Hybrid Search RAG
4. Multi-Index RAG
5. Query Expansion RAG
6. Adaptive RAG
7. Corrective RAG
8. Self Adaptive RAG
9. Hypothetical Document Embedding (HyDE)

Agents

LLM agents use large language models to handle complex queries, combining general and specialized knowledge, making them valuable across industries.

Multi-Agentic Frameworks – CrewAI, LangGraph



LLMs

Closed Source – OpenAI o1, GPT-4o, GPT-4o Mini, Gemini 1.5/1 Pro, Gemini 1.5 Flash-8B, Claude 3.5 Sonnet, Claude 3.5 Haiku etc.

Open LLMs – DeepSeek-v3, DeepSeek-R1, DeepSeek distilled models, Meta's Llama 3.2 1B, & 3B, Llama 3.1 405B, Microsoft's Phi3.5-Mini, Google's Gemma 2, Mistral 3B & 8B, Huggingface based Open-source Models etc.

Multimodal

Closed Source – GPT-4o, Gemini 1.5/1 Pro, Gemini 1.5 Flash-8B, Claude 3.5 Sonnet etc

Open Source – Meta's Llama 3.2 11B, 90B, Google's PaliGemma, LLaVA-1.5, Pixtral 12B, QwenVL, DeepSeek Janus-Pro 7B & 1 B, Huggingface based MultiModal

LLM Fine Tuning

1. Data Selection and Preprocessing
2. Transfer Learning and Domain Adaptation
3. Parameter-Efficient Fine-Tuning
4. Prompt Tuning and Instruction Fine-Tuning
5. Evaluation Metrics for Fine-Tuned Models
6. Safety and Bias Mitigation in Fine-Tuning
7. Hyperparameter Optimization

If you know all these, then you know most things in GenAI



Prompt Engineering

Prompt Techniques

1. Chain of Thought (CoT)
2. Few-Shot Chain of Thought (Few-Shot-CoT)
3. ReAct (Reasoning and Acting)
4. Tree of Thoughts (ToT)
5. Self-Consistency
6. Hypothetical Document Embeddings (HyDE)
7. Least-to-Most Prompting
8. Graph Prompting
9. Recursive Prompting
10. Generated Knowledge
11. Automatic Reasoning and Tool-Use (ART)
12. Automatic Prompt Engineer (APE)



Transformer Architecture

1. Self-Attention Mechanism
2. Positional Encoding
3. Multi-Head Attention:
4. Encoder-Decoder Architecture
5. Layer Normalization and Feed-Forward Layers
6. Pre-training and Fine-tuning
7. Scalability with Parallelization
8. Applications Across Domains



Cloud Support

Azure – Azure OpenAI, Azure AI Studio, Azure Machine Learning

AWS – Amazon Bedrock, Amazon Sagemaker, Amazon Q

Google – Vertex AI



LLM OPs

1. Model Deployment and Scaling
2. Monitoring and Logging
3. Versioning and Model Lifecycle Management
4. Cost Optimization
5. Feedback Loops and CI
6. Compliance and Security
7. Latency and Throughput Optimization
8. A/B Testing



LLM Evaluation

1. BLEU (Bilingual Evaluation Understudy)
2. ROUGE (Recall-Oriented Understudy for Gisting Evaluation)
3. Perplexity
4. Exact Match (EM)
5. Human Evaluation Scores
6. Bias and Fairness Metrics
7. Toxicity Scores



RAG

Vector DB - Pinecone, Weaviate, Qdrant, Chroma, Milvus, Vespa, LanceDB

Embedding Models – OpenAI's Embedding, Google's text-embedding-004, Huggingface Opensource Embedding Models

RAG Techniques

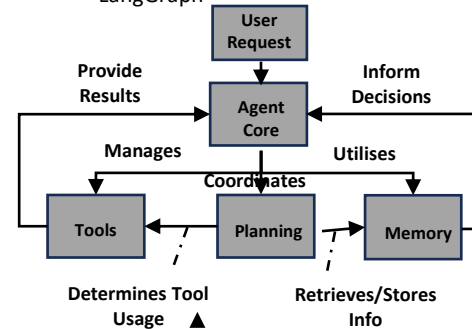
1. Basic RAG
2. Re-ranking RAG
3. Hybrid Search RAG
4. Multi-Index RAG
5. Query Expansion RAG
6. Adaptive RAG
7. Corrective RAG
8. Self Adaptive RAG
9. Hypothetical Document Embedding (HyDE)



Agents

LLM agents use large language models to handle complex queries, combining general and specialized knowledge, making them valuable across industries.

Multi-Agent Frameworks – CrewAI, LangGraph



LLMs

Closed Source – OpenAI o1, GPT-4o, GPT-4o Mini, Gemini 1.5/1 Pro, Gemini 1.5 Flash-8B, Claude 3.5 Sonnet, Claude 3.5 Haiku etc.

Open LLMs – DeepSeek-v3, DeepSeek-R1, DeepSeek distilled models, Meta's Llama 3.2 1B, & 3B, Llama 3.1 405B, Microsoft's Phi3.5-Mini, Google's Gemma 2, Ministral 3B & 8B, Huggingface based Opensource Models etc.



Multimodal

Closed Source – GPT-4o, Gemini 1.5/1 Pro, Gemini 1.5 Flash-8B, Claude 3.5 Sonnet etc

Open Source – Meta's Llama 3.2 11B, 90B, Google's PaliGemma, LLaVA-1.5, Pixtral 12B, QwenVL, DeepSeek Janus-Pro 7B & 1 B, Huggingface based MultiModal



LLM Fine Tuning

1. Data Selection and Preprocessing
2. Transfer Learning and Domain Adaptation
3. Parameter-Efficient Fine-Tuning
4. Prompt Tuning and Instruction Fine-Tuning
5. Evaluation Metrics for Fine-Tuned Models
6. Safety and Bias Mitigation in Fine-Tuning
7. Hyperparameter Optimization

LLM Frameworks (Across Layers) – LangChain/LlamaIndex